



POLITECNICO
MILANO 1863

Apache Kafka

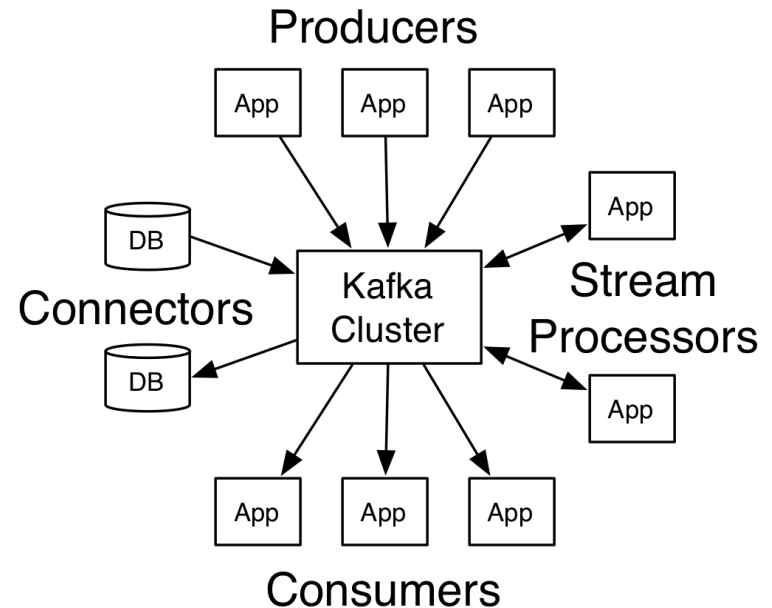
Alessandro Margara

`alessandro.margara@polimi.it`

`https://margara.faculty.polimi.it`

Event-based communication

- Architectural style to build distributed systems
- Components only interact by exchanging messages / events
 - Producers publish notifications of the events they observe
 - E.g., Security camera observes a person entering the room
 - Consumers subscribe to the events they are interested in
 - E.g., Security dashboard subscribes to all security-related events
- Producers and consumers need not know each other
 - The communication is mediated by a middleware service such as Kafka



Event-based communication

- Widely used to build large-scale dynamic systems, but why?
 - Space decoupling
 - Producers and consumers do not need to know each other
 - You can dynamically add new components without restarting/reconfiguring the system
 - E.g., a new sensor joins the network and produces new events
 - Synchronization decoupling
 - Producers are not blocked while producing events and consumers are notified asynchronously
 - Promotes scalability by removing explicit dependencies between producers and consumers
 - Time decoupling
 - If the event middleware can store events, producers and consumers need not be connected at the same time

Apache Kafka

- Apache Kafka is probably the most widely used platform for *event-based communication* today
- From the Website: “More than 80% of all Fortune 100 companies trust and use Kafka”
 - LinkedIn, Netflix, New York Times, Zalando, AirBnB, CISCO, Uber, PayPal, Spotify, Strava, Twitter, ...
- If you are interested in some use cases
 - <https://kafka.apache.org/uses>
 - <https://www.confluent.io/blog/category/use-cases/>
 - <https://www.confluent.io/designing-event-driven-systems/>

Apache Kafka

- Originally developed by LinkedIn to efficiently handle logs
- Open sourced in 2011
- Apache project since 2012
- In 2014, a group of developers who originally worked on Kafka at LinkedIn created a new company named Confluent
 - Develops Kafka
 - Offers a commercial platform that builds on Kafka
 - Additional management features
 - Solutions for Kafka-as-a-service
 - Since early 2021, Confluent is listed on Nasdaq


Apache Kafka overview

- Topic-based communication
 - Messages/events are put into Kafka queues/logs, each of them representing a topic
- Topics are persistent
 - Topics are stored on durable storage (disk) and replicated for fault-tolerance
- Consumers can read the same events more than once
 - Useful in the case of failure: a component can resume its state by replaying the history of events
- Scalability
 - Topics are partitioned to improve access performance and scalability

Case study


Event-based communication in microservices architectures

Case study: microservices architectures




- Architectural paradigm where an application is split into small independent services
- No shared-state
 - Each service has its own local state (database)
 - Services communicate only by RPC or by exchanging events

Case study: microservices architectures




- Orders processes orders from customers
- Orders are sent to the shipping service
- The shipping service queries the customer service for the address of the customer

Case study: microservices architectures




- The orders service invokes a command on the shipping service
 - Changes the state by adding a new order
- The shipping service queries the customer service

Case study: microservices architectures




- Focus on the interaction between the orders service and the shipping service
- We can use Kafka to mediate the communication
 - Instead of invoking a command on the shipping service ...
 - ... we register a new order event on the order topic

Case study: microservices architectures




- The two services are now decoupled!
- This brings several advantages

Case study: microservices architectures




- Orders does not need to know which services are consuming order events
 - We can change shipping service without Orders even noticing
 - We can add other services that consume order events without any change to the Orders service
 - E.g., a service that computes statistics about orders

Case study: microservices architectures




- The two services are decoupled in time
 - Shipping might be unavailable for a while ...
 - ... maybe because of maintenance or rollout of a new versions
 - When Shipping comes back online, it processes orders from the log
 - Orders can continue without noticing the temporary unavailability of Shipping

Case study: microservices architectures




- Now, let us focus on the interaction between Shipping and Customer
 - Shipping queries Customer for the address of customers before shipping an order
 - This is a synchronous query-response interaction ...
 - ... but synchronous interactions may negatively affect latency

Case study: microservices architectures




- Using an event-based paradigm, we can solve this problem by reversing the interaction between the services
 - Shipping does not query Customer
 - But it observes notifications of state (customer) changes from the Customer service

Case study: microservices architectures



- We can replicate the address of customers in the local database of the Shipping service
- When customers change their detail, the Customer service places a notification of change in the Customer updates topic
- Every interested service can read from Customer update
 - In particular, the Shipping service can update its local view of customers (address info)

Case study: microservices architectures



- Final design
- Benefits
 - Services are fully decoupled
 - No synchronous interaction
 - No problem in the case of unavailability of a service
 - Promotes evolution: add, remove, change services without other services even noticing
- Problem
 - Possible temporary inconsistency of Customer DB and Shipping DB

Case study: microservices architectures

- The above example introduces two key concepts
 - Event sourcing: events become the core element of the system
 - “Making events the source of truth” (rather than state)
 - “Turning the database inside-out”
 - Command query responsibility segregation (CQRS)
 - Separate the write path (commands) from the read path (queries)
 - Connects them through an asynchronous communication channel
 - An event log such as Apache Kafka